

# MNIST\_GAN\_Exercise

May 8, 2020

## 1 Generative Adversarial Network

In this notebook, we'll be building a generative adversarial network (GAN) trained on the MNIST dataset. From this, we'll be able to generate new handwritten digits!

GANs were [first reported on](#) in 2014 from Ian Goodfellow and others in Yoshua Bengio's lab. Since then, GANs have exploded in popularity. Here are a few examples to check out:

- [Pix2Pix](#)
- [CycleGAN & Pix2Pix in PyTorch, Jun-Yan Zhu](#)
- [A list of generative models](#)

The idea behind GANs is that you have two networks, a generator  $G$  and a discriminator  $D$ , competing against each other. The generator makes "fake" data to pass to the discriminator. The discriminator also sees real training data and predicts if the data it's received is real or fake. > \* The generator is trained to fool the discriminator, it wants to output data that looks *as close as possible* to real, training data. \* The discriminator is a classifier that is trained to figure out which data is real and which is fake.

What ends up happening is that the generator learns to make data that is indistinguishable from real data to the discriminator.

The general structure of a GAN is shown in the diagram above, using MNIST images as data. The latent sample is a random vector that the generator uses to construct its fake images. This is often called a **latent vector** and that vector space is called **latent space**. As the generator trains, it figures out how to map latent vectors to recognizable images that can fool the discriminator.

If you're interested in generating only new images, you can throw out the discriminator after training. In this notebook, I'll show you how to define and train these adversarial networks in PyTorch and generate new images!

```
In [1]: %matplotlib inline
```

```
import numpy as np
import torch
import matplotlib.pyplot as plt
```

```
In [2]: from torchvision import datasets
import torchvision.transforms as transforms

# number of subprocesses to use for data loading
num_workers = 0
```

```

# how many samples per batch to load
batch_size = 64

# convert data to torch.FloatTensor
transform = transforms.ToTensor()

# get the training datasets
train_data = datasets.MNIST(root='data', train=True,
                             download=True, transform=transform)

# prepare data loader
train_loader = torch.utils.data.DataLoader(train_data, batch_size=batch_size,
                                             num_workers=num_workers)

```

### 1.0.1 Visualize the data

In [3]: # obtain one batch of training images

```

dataiter = iter(train_loader)
images, labels = dataiter.next()
images = images.numpy()

```

```

# get one image from the batch
img = np.squeeze(images[0])


```

```

fig = plt.figure(figsize = (3,3))
ax = fig.add_subplot(111)
ax.imshow(img, cmap='gray')

```

Out[3]: <matplotlib.image.AxesImage at 0x7fb7d72f8588>



## 2 Define the Model

A GAN is comprised of two adversarial networks, a discriminator and a generator.

### 2.1 Discriminator

The discriminator network is going to be a pretty typical linear classifier. To make this network a universal function approximator, we'll need at least one hidden layer, and these hidden layers should have one key attribute: > All hidden layers will have a [Leaky ReLU](#) activation function applied to their outputs.

**Leaky ReLU** We should use a leaky ReLU to allow gradients to flow backwards through the layer unimpeded. A leaky ReLU is like a normal ReLU, except that there is a small non-zero output for negative input values.

**Sigmoid Output** We'll also take the approach of using a more numerically stable loss function on the outputs. Recall that we want the discriminator to output a value 0-1 indicating whether an image is *real* or *fake*. > We will ultimately use [BCEWithLogitsLoss](#), which combines a sigmoid activation function **and** binary cross entropy loss in one function.

So, our final output layer should not have any activation function applied to it.

```
In [4]: import torch.nn as nn
import torch.nn.functional as F

class Discriminator(nn.Module):

    def __init__(self, input_size, hidden_dim, output_size):
        super(Discriminator, self).__init__()

        # define hidden linear layers
        self.fc1 = nn.Linear(input_size, hidden_dim*4)
        self.fc2 = nn.Linear(hidden_dim*4, hidden_dim*2)
        self.fc3 = nn.Linear(hidden_dim*2, hidden_dim)

        # final fully-connected layer
        self.fc4 = nn.Linear(hidden_dim, output_size)

        # dropout layer
        self.dropout = nn.Dropout(0.3)

    def forward(self, x):
        # flatten image
        x = x.view(-1, 28*28)
        # all hidden layers
        x = F.leaky_relu(self.fc1(x), 0.2) # (input, negative_slope=0.2)
        x = self.dropout(x)
        x = F.leaky_relu(self.fc2(x), 0.2)
```

```

x = self.dropout(x)
x = F.leaky_relu(self.fc3(x), 0.2)
x = self.dropout(x)
# final layer
out = self.fc4(x)

return out

```

## 2.2 Generator

The generator network will be almost exactly the same as the discriminator network, except that we're applying a [tanh activation function](#) to our output layer.

**tanh Output** The generator has been found to perform the best with *tanh* for the generator output, which scales the output to be between -1 and 1, instead of 0 and 1.

Recall that we also want these outputs to be comparable to the *real* input pixel values, which are read in as normalized values between 0 and 1. > So, we'll also have to **scale our real input images to have pixel values between -1 and 1** when we train the discriminator.

I'll do this in the training loop, later on.

In [5]: `class Generator(nn.Module):`

```

def __init__(self, input_size, hidden_dim, output_size):
    super(Generator, self).__init__()

    # define hidden linear layers
    self.fc1 = nn.Linear(input_size, hidden_dim)
    self.fc2 = nn.Linear(hidden_dim, hidden_dim*2)
    self.fc3 = nn.Linear(hidden_dim*2, hidden_dim*4)

    # final fully-connected layer
    self.fc4 = nn.Linear(hidden_dim*4, output_size)

    # dropout layer
    self.dropout = nn.Dropout(0.3)

def forward(self, x):
    # all hidden layers
    x = F.leaky_relu(self.fc1(x), 0.2) # (input, negative_slope=0.2)
    x = self.dropout(x)
    x = F.leaky_relu(self.fc2(x), 0.2)
    x = self.dropout(x)
    x = F.leaky_relu(self.fc3(x), 0.2)
    x = self.dropout(x)
    # final layer with tanh applied
    out = F.tanh(self.fc4(x))

    return out

```

## 2.3 Model hyperparameters

```
In [6]: # Discriminator hyperparams

# Size of input image to discriminator (28*28)
input_size = 784
# Size of discriminator output (real or fake)
d_output_size = 1
# Size of last hidden layer in the discriminator
d_hidden_size = 32

# Generator hyperparams

# Size of latent vector to give to generator
z_size = 100
# Size of discriminator output (generated image)
g_output_size = 784
# Size of first hidden layer in the generator
g_hidden_size = 32
```

## 2.4 Build complete network

Now we're instantiating the discriminator and generator from the classes defined above. Make sure you've passed in the correct input arguments.

```
In [7]: # instantiate discriminator and generator
D = Discriminator(input_size, d_hidden_size, d_output_size)
G = Generator(z_size, g_hidden_size, g_output_size)

# check that they are as you expect
print(D)
print()
print(G)
```

```
Discriminator(
  (fc1): Linear(in_features=784, out_features=128, bias=True)
  (fc2): Linear(in_features=128, out_features=64, bias=True)
  (fc3): Linear(in_features=64, out_features=32, bias=True)
  (fc4): Linear(in_features=32, out_features=1, bias=True)
  (dropout): Dropout(p=0.3)
)
```

```
Generator(
  (fc1): Linear(in_features=100, out_features=32, bias=True)
  (fc2): Linear(in_features=32, out_features=64, bias=True)
  (fc3): Linear(in_features=64, out_features=128, bias=True)
  (fc4): Linear(in_features=128, out_features=784, bias=True)
  (dropout): Dropout(p=0.3)
```

)

---

## 2.5 Discriminator and Generator Losses

Now we need to calculate the losses.

### 2.5.1 Discriminator Losses

- For the discriminator, the total loss is the sum of the losses for real and fake images,  $d\_loss = d\_real\_loss + d\_fake\_loss$ .
- Remember that we want the discriminator to output 1 for real images and 0 for fake images, so we need to set up the losses to reflect that.

The losses will be binary cross entropy loss with logits, which we can get with `BCEWithLogitsLoss`. This combines a sigmoid activation function **and** binary cross entropy loss in one function.

For the real images, we want  $D(\text{real\_images}) = 1$ . That is, we want the discriminator to classify the real images with a label = 1, indicating that these are real. To help the discriminator generalize better, the labels are **reduced a bit from 1.0 to 0.9**. For this, we'll use the parameter `smooth`; if `True`, then we should smooth our labels. In PyTorch, this looks like `labels = torch.ones(size) * 0.9`

The discriminator loss for the fake data is similar. We want  $D(\text{fake\_images}) = 0$ , where the fake images are the *generator output*, `fake_images = G(z)`.

### 2.5.2 Generator Loss

The generator loss will look similar only with flipped labels. The generator's goal is to get  $D(\text{fake\_images}) = 1$ . In this case, the labels are **flipped** to represent that the generator is trying to fool the discriminator into thinking that the images it generates (fakes) are real!

```
In [8]: # Calculate losses
def real_loss(D_out, smooth=False):
    batch_size = D_out.size(0)
    # label smoothing
    if smooth:
        # smooth, real labels = 0.9
        labels = torch.ones(batch_size)*0.9
    else:
        labels = torch.ones(batch_size) # real labels = 1

    # numerically stable loss
    criterion = nn.BCEWithLogitsLoss()
    # calculate loss
    loss = criterion(D_out.squeeze(), labels)
    return loss
```

```
def fake_loss(D_out):
    batch_size = D_out.size(0)
    labels = torch.zeros(batch_size) # fake labels = 0
    criterion = nn.BCEWithLogitsLoss()
    # calculate loss
    loss = criterion(D_out.squeeze(), labels)
    return loss
```

## 2.6 Optimizers

We want to update the generator and discriminator variables separately. So, we'll define two separate Adam optimizers.

```
In [9]: import torch.optim as optim
```

```
# Optimizers
lr = 0.002

# Create optimizers for the discriminator and generator
d_optimizer = optim.Adam(D.parameters(), lr)
g_optimizer = optim.Adam(G.parameters(), lr)
```

---

## 2.7 Training

Training will involve alternating between training the discriminator and the generator. We'll use our functions `real_loss` and `fake_loss` to help us calculate the discriminator losses in all of the following cases.

### 2.7.1 Discriminator training

1. Compute the discriminator loss on real, training images
2. Generate fake images
3. Compute the discriminator loss on fake, generated images
4. Add up real and fake loss
5. Perform backpropagation + an optimization step to update the discriminator's weights

### 2.7.2 Generator training

1. Generate fake images
2. Compute the discriminator loss on fake images, using **flipped** labels!
3. Perform backpropagation + an optimization step to update the generator's weights

**Saving Samples** As we train, we'll also print out some loss statistics and save some generated "fake" samples.

```
In [10]: import pickle as pkl
```

```
# training hyperparams
num_epochs = 60

# keep track of loss and generated, "fake" samples
samples = []
losses = []

print_every = 400

# Get some fixed data for sampling. These are images that are held
# constant throughout training, and allow us to inspect the model's performance
sample_size=16
fixed_z = np.random.uniform(-1, 1, size=(sample_size, z_size))
fixed_z = torch.from_numpy(fixed_z).float()

# train the network
D.train()
G.train()
for epoch in range(num_epochs):

    for batch_i, (real_images, _) in enumerate(train_loader):

        batch_size = real_images.size(0)

        ## Important rescaling step ##
        real_images = real_images*2 - 1 # rescale input images from [0,1) to [-1, 1)

        # =====
        #                TRAIN THE DISCRIMINATOR
        # =====

        d_optimizer.zero_grad()

        # 1. Train with real images

        # Compute the discriminator losses on real images
        # smooth the real labels
        D_real = D(real_images)
        d_real_loss = real_loss(D_real, smooth=True)

        # 2. Train with fake images

        # Generate fake images
```



```

z = np.random.uniform(-1, 1, size=(batch_size, z_size))
z = torch.from_numpy(z).float()
fake_images = G(z)

# Compute the discriminator losses on fake images
D_fake = D(fake_images)
d_fake_loss = fake_loss(D_fake)

# add up loss and perform backprop
d_loss = d_real_loss + d_fake_loss
d_loss.backward()
d_optimizer.step()

# =====
#                               TRAIN THE GENERATOR
# =====
g_optimizer.zero_grad()

# 1. Train with fake images and flipped labels

# Generate fake images
z = np.random.uniform(-1, 1, size=(batch_size, z_size))
z = torch.from_numpy(z).float()
fake_images = G(z)

# Compute the discriminator losses on fake images
# using flipped labels!
D_fake = D(fake_images)
g_loss = real_loss(D_fake) # use real loss to flip labels

# perform backprop
g_loss.backward()
g_optimizer.step()

# Print some loss stats
if batch_i % print_every == 0:
    # print discriminator and generator loss
    print('Epoch [{:5d}/{:5d}] | d_loss: {:.64f} | g_loss: {:.64f}'.format(
        epoch+1, num_epochs, d_loss.item(), g_loss.item()))

## AFTER EACH EPOCH##
# append discriminator loss and generator loss
losses.append((d_loss.item(), g_loss.item()))

# generate and save sample, fake images
G.eval() # eval mode for generating samples

```

```

        samples_z = G(fixed_z)
        samples.append(samples_z)
        G.train() # back to train mode

# Save training generator samples
with open('train_samples.pkl', 'wb') as f:
    pickle.dump(samples, f)

Epoch [ 1/ 60] | d_loss: 1.3884 | g_loss: 0.7077
Epoch [ 1/ 60] | d_loss: 1.2285 | g_loss: 1.3171
Epoch [ 1/ 60] | d_loss: 1.3043 | g_loss: 0.7438
Epoch [ 2/ 60] | d_loss: 1.2051 | g_loss: 0.9774
Epoch [ 2/ 60] | d_loss: 0.9847 | g_loss: 1.7905
Epoch [ 2/ 60] | d_loss: 1.4263 | g_loss: 0.8025
Epoch [ 3/ 60] | d_loss: 1.2751 | g_loss: 0.8635
Epoch [ 3/ 60] | d_loss: 0.8772 | g_loss: 1.4054
Epoch [ 3/ 60] | d_loss: 1.0519 | g_loss: 1.1982
Epoch [ 4/ 60] | d_loss: 1.4932 | g_loss: 2.1488
Epoch [ 4/ 60] | d_loss: 1.0803 | g_loss: 1.3460
Epoch [ 4/ 60] | d_loss: 1.0941 | g_loss: 1.8086
Epoch [ 5/ 60] | d_loss: 1.0234 | g_loss: 1.6869
Epoch [ 5/ 60] | d_loss: 1.1167 | g_loss: 1.6820
Epoch [ 5/ 60] | d_loss: 1.4299 | g_loss: 0.8281
Epoch [ 6/ 60] | d_loss: 1.0371 | g_loss: 1.5915
Epoch [ 6/ 60] | d_loss: 1.1193 | g_loss: 1.3167
Epoch [ 6/ 60] | d_loss: 1.1302 | g_loss: 1.1695
Epoch [ 7/ 60] | d_loss: 1.3306 | g_loss: 1.0796
Epoch [ 7/ 60] | d_loss: 1.2251 | g_loss: 1.5026
Epoch [ 7/ 60] | d_loss: 1.0923 | g_loss: 1.2386
Epoch [ 8/ 60] | d_loss: 1.2108 | g_loss: 1.2196
Epoch [ 8/ 60] | d_loss: 1.2018 | g_loss: 1.0554
Epoch [ 8/ 60] | d_loss: 1.3157 | g_loss: 1.1538
Epoch [ 9/ 60] | d_loss: 1.2553 | g_loss: 0.9547
Epoch [ 9/ 60] | d_loss: 1.2158 | g_loss: 1.0619
Epoch [ 9/ 60] | d_loss: 1.1163 | g_loss: 1.1443
Epoch [10/ 60] | d_loss: 1.3878 | g_loss: 1.2848
Epoch [10/ 60] | d_loss: 1.4674 | g_loss: 0.8704
Epoch [10/ 60] | d_loss: 1.1789 | g_loss: 1.1959
Epoch [11/ 60] | d_loss: 1.1506 | g_loss: 1.1854
Epoch [11/ 60] | d_loss: 1.2284 | g_loss: 0.8909
Epoch [11/ 60] | d_loss: 1.2638 | g_loss: 0.9715
Epoch [12/ 60] | d_loss: 1.3462 | g_loss: 1.0427
Epoch [12/ 60] | d_loss: 1.2359 | g_loss: 1.2030
Epoch [12/ 60] | d_loss: 1.2173 | g_loss: 0.8969
Epoch [13/ 60] | d_loss: 1.1959 | g_loss: 1.4167
Epoch [13/ 60] | d_loss: 1.3602 | g_loss: 1.1672
Epoch [13/ 60] | d_loss: 1.3264 | g_loss: 0.9296

```

Epoch [	14/	60]	d_loss: 1.2738	g_loss: 0.9522
Epoch [	14/	60]	d_loss: 1.3238	g_loss: 1.0279
Epoch [	14/	60]	d_loss: 1.4493	g_loss: 0.8698
Epoch [	15/	60]	d_loss: 1.3362	g_loss: 0.9259
Epoch [	15/	60]	d_loss: 1.3352	g_loss: 1.0147
Epoch [	15/	60]	d_loss: 1.2065	g_loss: 1.2479
Epoch [	16/	60]	d_loss: 1.2397	g_loss: 1.6607
Epoch [	16/	60]	d_loss: 1.3028	g_loss: 1.0501
Epoch [	16/	60]	d_loss: 1.3963	g_loss: 1.0288
Epoch [	17/	60]	d_loss: 1.4047	g_loss: 0.8405
Epoch [	17/	60]	d_loss: 1.3119	g_loss: 0.8246
Epoch [	17/	60]	d_loss: 1.2459	g_loss: 0.9801
Epoch [	18/	60]	d_loss: 1.2128	g_loss: 1.1596
Epoch [	18/	60]	d_loss: 1.3366	g_loss: 0.8877
Epoch [	18/	60]	d_loss: 1.3857	g_loss: 0.9105
Epoch [	19/	60]	d_loss: 1.1782	g_loss: 1.4711
Epoch [	19/	60]	d_loss: 1.1180	g_loss: 1.7423
Epoch [	19/	60]	d_loss: 1.2733	g_loss: 0.8968
Epoch [	20/	60]	d_loss: 1.1236	g_loss: 1.2875
Epoch [	20/	60]	d_loss: 1.2618	g_loss: 1.0150
Epoch [	20/	60]	d_loss: 1.3100	g_loss: 0.8772
Epoch [	21/	60]	d_loss: 1.3610	g_loss: 0.8693
Epoch [	21/	60]	d_loss: 1.2332	g_loss: 0.9685
Epoch [	21/	60]	d_loss: 1.2803	g_loss: 1.3458
Epoch [	22/	60]	d_loss: 1.3346	g_loss: 0.9456
Epoch [	22/	60]	d_loss: 1.2718	g_loss: 1.0138
Epoch [	22/	60]	d_loss: 1.3024	g_loss: 0.9630
Epoch [	23/	60]	d_loss: 1.3203	g_loss: 1.2992
Epoch [	23/	60]	d_loss: 1.3566	g_loss: 0.9330
Epoch [	23/	60]	d_loss: 1.2482	g_loss: 1.1637
Epoch [	24/	60]	d_loss: 1.4943	g_loss: 1.0669
Epoch [	24/	60]	d_loss: 1.2290	g_loss: 1.0777
Epoch [	24/	60]	d_loss: 1.3320	g_loss: 0.9260
Epoch [	25/	60]	d_loss: 1.2518	g_loss: 1.0554
Epoch [	25/	60]	d_loss: 1.3245	g_loss: 0.8181
Epoch [	25/	60]	d_loss: 1.3363	g_loss: 1.2943
Epoch [	26/	60]	d_loss: 1.1900	g_loss: 1.0788
Epoch [	26/	60]	d_loss: 1.2416	g_loss: 0.9704
Epoch [	26/	60]	d_loss: 1.4097	g_loss: 0.9517
Epoch [	27/	60]	d_loss: 1.5551	g_loss: 0.8930
Epoch [	27/	60]	d_loss: 1.2183	g_loss: 1.0185
Epoch [	27/	60]	d_loss: 1.3388	g_loss: 1.0019
Epoch [	28/	60]	d_loss: 1.3018	g_loss: 1.3341
Epoch [	28/	60]	d_loss: 1.2652	g_loss: 1.0707
Epoch [	28/	60]	d_loss: 1.3204	g_loss: 0.8523
Epoch [	29/	60]	d_loss: 1.2488	g_loss: 1.1985
Epoch [	29/	60]	d_loss: 1.2772	g_loss: 0.9125
Epoch [	29/	60]	d_loss: 1.2434	g_loss: 0.9989

Epoch [	30/	60]	d_loss: 1.3071	g_loss: 1.1866
Epoch [	30/	60]	d_loss: 1.2758	g_loss: 0.9508
Epoch [	30/	60]	d_loss: 1.3255	g_loss: 1.0869
Epoch [	31/	60]	d_loss: 1.2181	g_loss: 0.9239
Epoch [	31/	60]	d_loss: 1.2393	g_loss: 0.9209
Epoch [	31/	60]	d_loss: 1.3457	g_loss: 1.0523
Epoch [	32/	60]	d_loss: 1.2711	g_loss: 1.1656
Epoch [	32/	60]	d_loss: 1.3096	g_loss: 1.0965
Epoch [	32/	60]	d_loss: 1.3004	g_loss: 1.0638
Epoch [	33/	60]	d_loss: 1.2859	g_loss: 1.0291
Epoch [	33/	60]	d_loss: 1.2190	g_loss: 1.0418
Epoch [	33/	60]	d_loss: 1.3476	g_loss: 0.9261
Epoch [	34/	60]	d_loss: 1.3933	g_loss: 0.9802
Epoch [	34/	60]	d_loss: 1.2783	g_loss: 0.9696
Epoch [	34/	60]	d_loss: 1.4638	g_loss: 1.3835
Epoch [	35/	60]	d_loss: 1.4258	g_loss: 2.0603
Epoch [	35/	60]	d_loss: 1.2850	g_loss: 0.9085
Epoch [	35/	60]	d_loss: 1.4509	g_loss: 0.9419
Epoch [	36/	60]	d_loss: 1.3304	g_loss: 0.9733
Epoch [	36/	60]	d_loss: 1.2524	g_loss: 0.8699
Epoch [	36/	60]	d_loss: 1.3029	g_loss: 0.9594
Epoch [	37/	60]	d_loss: 1.3028	g_loss: 0.8500
Epoch [	37/	60]	d_loss: 1.3430	g_loss: 0.8029
Epoch [	37/	60]	d_loss: 1.3396	g_loss: 0.9733
Epoch [	38/	60]	d_loss: 1.3225	g_loss: 1.0202
Epoch [	38/	60]	d_loss: 1.3669	g_loss: 0.8398
Epoch [	38/	60]	d_loss: 1.4268	g_loss: 0.9517
Epoch [	39/	60]	d_loss: 1.3312	g_loss: 0.9041
Epoch [	39/	60]	d_loss: 1.2339	g_loss: 0.9726
Epoch [	39/	60]	d_loss: 1.3242	g_loss: 0.9406
Epoch [	40/	60]	d_loss: 1.3856	g_loss: 0.8984
Epoch [	40/	60]	d_loss: 1.2799	g_loss: 0.9118
Epoch [	40/	60]	d_loss: 1.2946	g_loss: 1.2216
Epoch [	41/	60]	d_loss: 1.3181	g_loss: 0.9175
Epoch [	41/	60]	d_loss: 1.2278	g_loss: 0.8886
Epoch [	41/	60]	d_loss: 1.2813	g_loss: 0.9294
Epoch [	42/	60]	d_loss: 1.2732	g_loss: 0.9880
Epoch [	42/	60]	d_loss: 1.3195	g_loss: 0.9519
Epoch [	42/	60]	d_loss: 1.3454	g_loss: 1.0715
Epoch [	43/	60]	d_loss: 1.2020	g_loss: 1.6028
Epoch [	43/	60]	d_loss: 1.3432	g_loss: 0.8819
Epoch [	43/	60]	d_loss: 1.3206	g_loss: 0.9768
Epoch [	44/	60]	d_loss: 1.3043	g_loss: 0.8567
Epoch [	44/	60]	d_loss: 1.3472	g_loss: 0.9475
Epoch [	44/	60]	d_loss: 1.3753	g_loss: 0.8114
Epoch [	45/	60]	d_loss: 1.2313	g_loss: 1.1159
Epoch [	45/	60]	d_loss: 1.2942	g_loss: 0.8438
Epoch [	45/	60]	d_loss: 1.2299	g_loss: 0.9296

Epoch [	46/	60]	d_loss: 1.2955	g_loss: 1.0312
Epoch [	46/	60]	d_loss: 1.2527	g_loss: 1.2390
Epoch [	46/	60]	d_loss: 1.3621	g_loss: 0.8361
Epoch [	47/	60]	d_loss: 1.2230	g_loss: 1.0662
Epoch [	47/	60]	d_loss: 1.2319	g_loss: 0.9217
Epoch [	47/	60]	d_loss: 1.3416	g_loss: 0.9799
Epoch [	48/	60]	d_loss: 1.2978	g_loss: 1.0461
Epoch [	48/	60]	d_loss: 1.2277	g_loss: 1.1099
Epoch [	48/	60]	d_loss: 1.2239	g_loss: 1.4972
Epoch [	49/	60]	d_loss: 1.3298	g_loss: 0.9697
Epoch [	49/	60]	d_loss: 1.2926	g_loss: 1.0085
Epoch [	49/	60]	d_loss: 1.3520	g_loss: 0.8952
Epoch [	50/	60]	d_loss: 1.1682	g_loss: 0.9983
Epoch [	50/	60]	d_loss: 1.1878	g_loss: 0.9267
Epoch [	50/	60]	d_loss: 1.4014	g_loss: 0.9846
Epoch [	51/	60]	d_loss: 1.3085	g_loss: 0.9878
Epoch [	51/	60]	d_loss: 1.2962	g_loss: 0.9604
Epoch [	51/	60]	d_loss: 1.3530	g_loss: 0.9212
Epoch [	52/	60]	d_loss: 1.3382	g_loss: 0.8317
Epoch [	52/	60]	d_loss: 1.2786	g_loss: 0.9968
Epoch [	52/	60]	d_loss: 1.3102	g_loss: 0.9271
Epoch [	53/	60]	d_loss: 1.2568	g_loss: 0.8155
Epoch [	53/	60]	d_loss: 1.2866	g_loss: 1.0641
Epoch [	53/	60]	d_loss: 1.3421	g_loss: 0.9050
Epoch [	54/	60]	d_loss: 1.3491	g_loss: 1.0781
Epoch [	54/	60]	d_loss: 1.2876	g_loss: 0.8498
Epoch [	54/	60]	d_loss: 1.2902	g_loss: 0.9837
Epoch [	55/	60]	d_loss: 1.2502	g_loss: 1.1737
Epoch [	55/	60]	d_loss: 1.2667	g_loss: 1.0283
Epoch [	55/	60]	d_loss: 1.3517	g_loss: 0.8530
Epoch [	56/	60]	d_loss: 1.2095	g_loss: 0.9883
Epoch [	56/	60]	d_loss: 1.2448	g_loss: 0.9131
Epoch [	56/	60]	d_loss: 1.2924	g_loss: 0.9068
Epoch [	57/	60]	d_loss: 1.2963	g_loss: 1.1111
Epoch [	57/	60]	d_loss: 1.2914	g_loss: 0.8530
Epoch [	57/	60]	d_loss: 1.4525	g_loss: 0.8508
Epoch [	58/	60]	d_loss: 1.3299	g_loss: 0.8455
Epoch [	58/	60]	d_loss: 1.2705	g_loss: 1.1248
Epoch [	58/	60]	d_loss: 1.3951	g_loss: 1.0257
Epoch [	59/	60]	d_loss: 1.3116	g_loss: 0.9647
Epoch [	59/	60]	d_loss: 1.2430	g_loss: 0.9689
Epoch [	59/	60]	d_loss: 1.2870	g_loss: 1.1637
Epoch [	60/	60]	d_loss: 1.3487	g_loss: 1.2861
Epoch [	60/	60]	d_loss: 1.2684	g_loss: 1.0935
Epoch [	60/	60]	d_loss: 1.2567	g_loss: 1.0948

## 2.8 Training loss

Here we'll plot the training losses for the generator and discriminator, recorded after each epoch.

```
In [11]: fig, ax = plt.subplots()
         losses = np.array(losses)
         plt.plot(losses.T[0], label='Discriminator')
         plt.plot(losses.T[1], label='Generator')
         plt.title("Training Losses")
         plt.legend()
```

```
Out[11]: <matplotlib.legend.Legend at 0x7fb7cda36f28>
```



## 2.9 Generator samples from training


Here we can view samples of images from the generator. First we'll look at the images we saved during training.

```
In [12]: # helper function for viewing a list of passed in sample images
         def view_samples(epoch, samples):
             fig, axes = plt.subplots(figsize=(7,7), nrows=4, ncols=4, sharey=True, sharex=True)
             for ax, img in zip(axes.flatten(), samples[epoch]):
                 img = img.detach()
                 ax.xaxis.set_visible(False)
                 ax.yaxis.set_visible(False)
                 im = ax.imshow(img.reshape((28,28)), cmap='Greys_r')
```

```
In [13]: # Load samples from generator, taken while training
         with open('train_samples.pkl', 'rb') as f:
             samples = pickle.load(f)
```

These are samples from the final training epoch. You can see the generator is able to reproduce numbers like 1, 7, 3, 2. Since this is just a sample, it isn't representative of the full range of images this generator can make.

```
In [14]: # -1 indicates final epoch's samples (the last in the list)
         view_samples(-1, samples)
```




Below I'm showing the generated images as the network was training, every 10 epochs.

```
In [15]: rows = 6 # split epochs into 10, so 100/10 = every 10 epochs
         cols = 6
         fig, axes = plt.subplots(figsize=(7,12), nrows=rows, ncols=cols, sharex=True, sharey=True)
```

```
for sample, ax_row in zip(samples[:, :int(len(samples)/rows)], axes):
    for img, ax in zip(sample[:, :int(len(sample)/cols)], ax_row):
        img = img.detach()
        ax.imshow(img.reshape((28,28)), cmap='Greys_r')
        ax.xaxis.set_visible(False)
        ax.yaxis.set_visible(False)
```





It starts out as all noise. Then it learns to make only the center white and the rest black. You can start to see some number like structures appear out of the noise like 1s and 9s.


## 2.10 Sampling from the generator

We can also get completely new images from the generator by using the checkpoint we saved after training. **We just need to pass in a new latent vector  $z$  and we'll get new samples!**

```
In [16]: # randomly generated, new latent vectors
sample_size=16
rand_z = np.random.uniform(-1, 1, size=(sample_size, z_size))
rand_z = torch.from_numpy(rand_z).float()

G.eval() # eval mode
# generated samples
rand_images = G(rand_z)

# 0 indicates the first set of samples in the passed in list
# and we only have one batch of samples, here
view_samples(0, [rand_images])
```



In [ ]: